



ADAPTIVE POKER BOT BASED ON PLAYER-MODELLING

Bjørn Hasager Vinther & Nicolai Guldbæk Holst

bhas@itu.dk

ngul@itu.dk

A thesis submitted for the degree of

Bachelor Softwaredevelopment

May 2014

Abstract

Nothing yet...

Texas hold'em poker

Poker is arguably one of the most popular gambling games. It combines psychology, statistics, and luck into a single game with simple rules. Even though the rules of poker are simple, the game itself is hard to master. It is played with a standard deck of 52 cards. This thesis only concerns Texas Hold'em limit poker. Therefore, the basic rules and the line of play is explained below. [?].

Player actions

- Fold** The player gives up and the player is out until the start of the next round. This action is always possible.
- Check** The player does not bet any chips. This action is only possible if no other player has placed a bet.
- Bet** The player bets an amount equal to the blind. All opponents have to match his bid in order to stay in the game.
- Call** The player place a bet that matches the highest bid of the opponents.
- Raise** The player place bet higher than the bet of the opponents. All opponents have to match this bid in order to stay in the game. This action is only possible if another player has placed a bet.
- All-in** The player bets all his chips. The player is still in the game but he will not be able to act any more. This action is only possible if you cannot afford to call, bet or raise.

Keywords

Dealer	The players takes turn being the dealer in the beginning of each round. The player to the left of the dealer is always the first to take action in the bidding round.
Blind	A fixed amount that the two players has to pay in the start of each round.
Hole cards	A pair of private cards that is dealt to each player.
Community cards	The five public cards that are dealt to the table and are shared between all players.
Bidding round	In the bidding round the players takes turn performing an action. A bidding round will occur in every game state.
Action	An action can be performed when it is the players turn to act during the bidding round.
Pot	The sum of all the bids that have been placed in the round. The winner of the round wins the pot.
Chips	The amount of money a player has.
Round	A round consists of 4 game states: pre-flop, flop, turn, and river. In the beginning of each round the blind are paid.
Pre-flop	The first game state. In this state the hole cards are dealt.
Flop	The second game state. In this state the first three community cards are dealt.
Turn	The third game state. In this state the fourth community card is dealt.
River	The final game state. In this state the fifth community card is dealt. After the bidding round a showdown will take place.
Showdown	The phase where all players who have not folded show their hole cards and the winner is determined.
Hand	The best combination of the players hole cards and the community cards. See table 1.

Game play

A Texas hold'em poker game consists of multiple rounds. A round consists of four game states: pre-flop, flop, turn and river. Each game state starts with cards being dealt and then the bidding begins. In the preflop state the dealer will deal 2 cards to each of the players this is their hole cards. In the bidding the players will take turn performing their actions. If only one player is left after the bidding that player wins the pot, otherwise the game continues to the next game state. If multiple players are still left after the bidding a showdown will start. During showdown each player left will reveal their hole cards and a winner will be found. The winner wins the pot. In case of a draw the pot is split between the winners.

The amount of chips is the main indicator of how well the player is doing. The main goal for the players is to increase their amount of chips.

The bidding

The bidding is where the players in turn perform their actions. The actions include: call, bet, raise, fold, check, all-in. The player to the left of the dealer is always the first to act.

A player can play aggressively by betting or raising which will increase the cost for the other players. Likewise a player can also play defensively by calling or checking which will not increase the cost for other players. If a player decides to fold he will lose what he betted that round. Whenever a player chooses to play aggressively all other players must either fold or call in order for the bidding round to stop. The bidding continues until all players have called the aggressor or folded.

Rules for determining the winner

The rules for finding the best hand are quite simple. Each player has to create the best possible hand they can using the hole cards and the community cards. We refer to the rank of a card (2-A) as the card rank and the rank of the hand (see table 1) as the rank. First the rank of the hand is found. The order of the cards is insignificant. In case two players have the same rank the winner will be determined by the card ranks of the given hand. For instance a pair of kings is better than a pair of jacks. No card suit is better than another.

Starting from the top (the best hand) in table 1 we have the explanations:
High card is the highest card.
One pairs is having two cards with the same card rank.
Two pairs is having two pairs
Three of a kind is having three cards with the same card rank
Straight is having five cards in a row (e.g 10-A or 4-9)
Flush is having five cards with the same suit
Full house is having three of a kind and a pair
Four of a kind is having four cards with the same card rank
Straight Flush is the same as a royal flush except there is no requirement for the highest card.
Royal Flush is having a straight all in the same suit. Furthermore the highest card have to be an ace.

rank	name	example hand
1	Royal flush	A♣ K♣ Q♣ J♣ 10♣
2	Straight flush	7♣ 6♣ 5♣ 4♣ 3♣
3	Four of a kind	K♣ K♠ K♦ K♥ 10♣
4	Full house	K♣ K♠ K♦ Q♥ Q♣
5	Flush	K♥ Q♥ 5♥ 3♥ 2♥
6	Straight	A♣ K♠ Q♦ J♥ 10♣
7	Three of a kind	A♣ A♦ A♠ J♣ 10♣
8	Two pairs	A♣ A♦ 5♠ 5♣ 4♣
9	One pair	A♣ A♥ J♣ 9♠ 2♥
10	High card	A♣ K♦ 6♥ 5♥ 3♥

Table 1: Ranks of different hands in Texas hold'em poker sorted best to worst.

Preface

This bachelor thesis consists of 15 ETCS points. The thesis has been written exclusively by Nicolai Guldbæk Holst and Bjørn Hasager Vinther, both studying *Software development* at the IT University of Copenhagen. It has been written in the spring semester 2015. Our supervisor was Kasper Støy.

Introduction

Poker is arguably the most popular card game in the world. It involves intelligence, psychology, and luck. Poker has a lot of hidden information which requires the players to make decisions based on qualified guesses.

Reading a player refers to the process of figuring out the players strategy and is a major element of poker. The top players are able to read their opponents and adapt to their strategy in order to gain an advantage. Most rounds are won before the showdown simply by one player outplaying the others.

Another big element of poker is statistics. Since the players do not know the cards that will be dealt throughout the game, players must calculate the likelihood of their hand ending up being the winning hand. The more likely the player is to win the more aggressive he can play to maximize his profit.

In the field of artificial intelligence, games are interesting because of their well-defined game rules and success criteria. Computers have already mastered some of the popular games, one example is the chess computer Deep Blue which won against Garry Kasparov, the world champion of chess at the time. Chess is a game of perfect information, as no information is hidden from the players. Since then the interest of the artificial intelligence research has shifted towards games with imperfect information. These types of games presents new challenges such as deception and hidden information. Poker is an example of a game with imperfect information.

Developing an algorithm capable of playing poker is not only limited to the domain of poker but can end up having a future applications in other domains as well. In essence a game presents a challenge for the players to solve. How the player approaches this challenge and what strategy the player uses to solve it, are what determines the players success.

“Bowling says the findings in this new research are especially valuable because they give us a hint at the scale of problems AI can solve. ... Solving a game as complex as heads-up limit Texas hold ’em could mean a breakthrough in our conception of how big is too big”

- Lily Hay Newman, *Using AI to Study Poker Is Really About Solving Some of the World’s Biggest Problems* line 46-52

In may 2015 a contest was held with four of the best professional poker players in the world. Each of the players had to play 20.000 hands of heads-up no-limit Texas Hold’em against the currently best poker bot in the world Claudico. Claudico was able to end up with more chips than one of the four players which proves that artificial intelligence in regards to poker have come a long way.

The goal of this thesis is to see if it is possible to train a computer to play poker and adapt to its opponents strategy. In order to achieve this goal we first need to solve the following problem statements:

Problem statements

1. How can one determine the strength of a poker hand?
2. How can we develop a default strategy without having information about the opponents?
3. How can one develop a strategy adapted to the opponents.

Our thesis is divided into three chapters each of which focuses on one of the three problem statements.

In chapter 1 1 we find a solution to problem statement one. We develop a subsystem which is able to estimate the probability of winning for any set of hole cards in any poker state. The subsystem can calculate the probability of winning with an error percentage of one percent and it takes less than a second on average.

In chapter 2 2 ...

In chapter 3 3 ...

Contents

1	Determine the strength of a poker hand	10
1.1	Design	10
1.1.1	Monte Carlo method	12
1.2	Test	12
1.2.1	Number of simulations	12
1.2.2	Accuracy of the calculator	14
1.3	Discussion	16
1.4	Conclusion	17
2	Learning a default strategy	18
2.1	Design	18
2.1.1	Artificial neural network (ANN)	19
2.1.2	First ANN design	21
2.1.3	Second ANN design	23
2.1.4	Third ANN design	24
2.2	Test	24
2.2.1	Find the TNE of the first ANN	25
2.2.2	Find the TNE of the second ANN	26
2.2.3	Find the TNE of the third ANN	27
2.3	Discussion	28
2.4	Conclusion	29
3	Learning to adapt to opponents	30
3.0.1	Player modeling	30
3.1	How can we model a player?	31
3.2	How can we model a player dynamically?	32
3.3	Test	32
3.4	Discussion	32
3.5	Conclusion	32
4	Discussion	33
5	Conclusion	34
A		
	Glossary	39

1 Determine the strength of a poker hand

Our first step towards developing an adaptive poker bot is to find a way to determine the strength of any given hand in any game state. In this section we will answer the following problem statement:

Problem statement 1

How can one determine the strength of a poker hand?

The strength of a hand reflects the probability of winning so the stronger a hand is the more likely one are to win. Since the player does not know the outcome of the community cards during a round of poker, the player has to calculate the probability of winning based on the possible outcomes of the community cards. It is too time consuming to check every outcome as there are more than 250 millions different outcomes of community cards alone. In order to find the probability of winning without having to check all possible outcomes, one can instead make an estimate rather than calculating the true probability.

1.1 Design

When trying to estimate the strength of a hand, two options exist.

The first option is to create a simplified formula. Such formulas already exist, but since they are very simple they tend to be rather inaccurate. This option is straight forward, but the disadvantage is, that it is hard to make a formula that is accurate for every game state.

The second option is to use the Monte Carlo method to simulate a large amount of games and get the distribution of outcomes. This distribution can then be used to find the probability for any of the outcomes.

For a human player a simplified formula is necessary but due to the computational power of a computer the Monte Carlo method is optimal for the APC. The computer can perform thousands of simulations in no time. This method also gives a trade-off between accuracy and the number of simulations. This allows one to adjust the accuracy of the probability by adjusting the number of simulations. The major poker sites also use the Monte Carlo method to determine each players probability of winning.

The solution is implemented as a subsystem that uses the Monte Carlo method. We will refer to this subsystem as the calculator.

Algorithm 1: Pseudo-code for a single simulation

	Data: Hole cards, number of opponents, community cards (optional)
	Result: win = 1, draw = 0, lose = -1
1	Random missing community cards;
2	Determine rank of players hand;
3	foreach <i>opponent</i> do
4	Determine rank of opponents hand;
5	if <i>opponent has stronger hand</i> then
6	return -1;
7	end
8	else if <i>opponent has same hand</i> then
9	return 0;
10	end
11	end
12	return 1;

The calculator takes three arguments: the hole cards of the player, the number of opponents, and the community cards (optional). The calculator simulates a pre-defined number of rounds and returns an object containing the distribution of wins, draws, and loses.

The calculator considers it a win only if the hand beats every other hand of the opponents.

In order to ensure the quality of the calculator, we have the following requirements:

- It must be able to return the probability of winning with a given hand in any game state with up to ten players.
- It must have a maximum error percentage of one percent. (deviation from the true probability)
- It must calculate the probability in less than three seconds.

For requirement one we chose ten players because ten is the maximum number of players in a single round. For requirement two we chose an error percentage of one because we found it adequate. For requirement three we

chose three second as it seemed fast enough since most players usually have 20-30 seconds for each decision making.

1.1.1 Monte Carlo method

The Monte Carlo method is used to find the distribution of outcomes for a domain. Given a set of user defined inputs the Monte Carlo method performs the simulation to find the outcome. For each simulation the method tracks the outcomes and as the number of simulations increases, so will the accuracy of the distribution. This distribution can help to get a better understanding of the domain.

1.2 Test

As mentioned earlier the number of simulations affects the accuracy of the result. First we test what number of simulations is needed to fulfil the requirements. In order to find the accuracy of the probabilities we first need to know the true probabilities. For this we use caniwini ?? . Caniwini is a website that has simulated all possible outcomes of any pre-flop hands. Caniwini's results are limited to the pre-flop game state, so we can only test the calculator for this game state.

After we find an optimal number of simulations we will test the accuracy of the calculator in different game states using this number of simulations.

1.2.1 Number of simulations

To find the optimal number of simulations for the calculator we test a single hand with a different number of simulations. The calculator calculates the probability 50 times for each test. Each test is performed with the hand $J\clubsuit J\heartsuit$ in pre-flop with one opponent and the true probability of this is 77,1 %. We then find the absolute max error and the average computation time of all calculations. The error percentage in table 2 is found using the formula:

$$err = P_O - P_T$$

Here P_O is the probability found by the calculator and P_T is the true probability.

We plotted the results of each test in a graph, see figure 1, 2, and 3. Each test result is indicated with a red dot. The three graphs clearly shows that when using a higher amount of simulations, the more accurate the probability becomes.

The error percentage can be seen in table 5.




Figure 1: Result of the calculator with 1000 simulations




Figure 2: Result of the calculator with 10.000 simulations




Figure 3: Result of the calculator with 50.000 simulations

We settled for 50.000 as the number of simulations for our calculator. This satisfies our requirements.

1.2.2 Accuracy of the calculator

To test if the calculator can calculate the correct probabilities we find the probability of a number of different pre-flop scenarios and compare the result with the results by caniwins. Every test have been preformed with 50.000 simulations against one opponent. The result can be seen in table 2. From the results we can see that for all hole cards the error percentage is less than one.

hole cards	P_O (%)	P_T (%)	error (%)
A♣ A♦	85,2	84,9	0,3
8♣ 8♦	67,9	68,7	0,8
Q♣ k♣	62,8	62,4	0,4
A♥ 8♠	58,8	60,5	0,7
J♠ Q♦	57,2	56,9	0,3
10♥ J♥	56,7	56,2	0,5
3♦ 3♠	53,0	52,8	0,2
2♦ 2♥	49,5	49,4	0,1
9♦ 3♠	37,8	37,4	0,4
2♦ 7♦	35,5	35,4	0,1
2♦ 7♥	31,9	31,7	0,2

Table 2: Test results for different hole cards in pre-flop with one opponent.

The pair of fives in table 3 have a zero percent chance of winning at the river, because the monte carlo does not count a draw as a win.

hole cards	community cards	pre-flop (%)	flop (%)	turn (%)	river (%)
A♣ K♦	3♠ 7♣ T♣ 8♦ 2♠	64,5	54,4	41,7	35,6
2♦ 5♠	3♠ 6♦ K♦ 4♥ 7♠	29,7	28,2	92,8	94,2
8♠ 9♠	2♦ K♠ 3♠ A♣ T♠	49,1	53,5	38,5	99,7
5♠ 5♣	A♦ K♥ Q♠ J♦ T♥	58,5	48,0	34,4	0,0

Table 3: Test results for different hole cards in pre-flop, flop, turn, and river with one opponent

Despite a pair of aces is a good hand, the number of opponents makes the chances of winning a lot less

opponents	1	2	3	4	5	6	7	8	9
win probability (%)	85,2	73,4	64,0	56,1	49,5	43,9	39,1	35,0	31,6

Table 4: Chance of winning with the hole cards A♠ A♣ in pre-flop

In table 5 you can see the combined result. The range is the difference between the lowest and highest result and the max error is the maximum deviation from caniwins.

The time have been calculated by finding an average of 100 calculations of the same hole cards.

simulations	range (%)	max error (%)	time (seconds)
1000	5,2	2,7	~ 0,01
10.000	2,2	1,5	~ 0.04
50.000	0,9	0,7	~ 0.15

Table 5: Combined test results from running the calculator with different numbers of simulations.

1.3 Discussion

For the implementation of the calculator the Monte Carlo method was. This method works well for the needs. The calculator can calculate the probability with an error percentage of less than one percent when using 50.000 simulations. This solution can be used for any poker state with up to ten players. At first it did not seem necessary to optimise the calculator any further, as the running time is acceptable. But in regards to later use or if were to scale the project, we decided to make an obvious optimisation by making it multi-threaded which speeded up the calculations by roughly 3 times.

Since we compare the results of the calculator to caniwins, one would have to ensure the reliability of that source. It is not directly shown who has written the article but when trying to contact the owner of the write we are directed to Bobby. It is very difficult to know if Bobby is the author and whether or not he has any reputation and expertise in the field. However after finding out that the probability from caniwins heads-up and 10 players is very close to the results from the calculator we concluded that caniwins were a reliable source of information.

Alternatively we could have created our own formula to calculate a rank, or used an existing one, for instance the Chen formula. By using this method we will get a less accurate result but it will be easier to calculate. We would also have to create a formula for each game state which would cause even more work. Since performance is not a problem for our calculator we chose to use the Monte Carlo method.

1.4 Conclusion

In this section we have answered the question:

Problem statement 1

How can we predict the probability of ending up with the winning hand?

We have implemented a subsystem called the calculator that can estimate the probability of winning with a set of hole cards. The calculator uses the Monte Carlo method and it works for every poker state with up to ten players. We have found that 50.000 simulations is a good number of simulations for our calculator.

We compare our results to caniwins, which is a website that calculates the actual probabilities of all the 169 combinations of hole cards. The calculator has a maximal error percentage of one percent and performs the simulation in less than a second.

2 Learning a default strategy

In the previous chapter we created a calculator that can determine the strength of a hand by calculating the probability of winning. In this chapter we shall use the calculator to calculate the strength of the hands.

When the APC first joins a poker game it has no information about the opponent, and in this case it must use a default strategy while it gathers more information.

In this chapter we will find a solution to the problem statement:

Problem statement 2

How can one develop a default strategy without having information about the opponents?

Even though players have different strategies they often have some decisions in common. Most players tend to play more aggressively the better their chances are of winning and likewise most players will fold if they have a weak hand. These tendencies can be used in a default strategy. The popular decisions are more likely to be good. For instance, most players agree that it is unwise to fold a pair of aces in the pre-flop.

The default strategy has to work against every strategy, therefore it is impossible for it to be better than all of them. The goal for the default strategy is not to win, although that is preferable, but instead to reduce the losses while it gathers information about the opponent.

2.1 Design

To develop a strategy in poker the two most commonly used options are to either directly program the procedures in the code or to create a self-learning algorithm.

Programming the procedures in the code requires the programmer to have a deep insight in how to play poker and how to make the optimal decisions during a game. One can also use the expertise of professional poker players in case one lacks the insight.

The self-learning algorithm uses the concept watch and learn by observing

other players and trying to learn the strategy behind their decisions. This method requires that the algorithm has someone to observe.

Since we do not have any particular insight in how to play poker and do not have expertise from any professional poker players, we will implement a self-learning algorithm. Additionally, by using this method the computer is not limited by our understanding of the game. We will use data from real life poker games for the algorithm. The data is further described in section 2.2.

To implement the self-learning algorithm we use an artificial neural network (ANN), see section 2.1.1. The ANN is well suited for finding patterns of the players decisions.

Our goal is to design an ANN with a total network error (TNE) of five percent or less. We find five percent to be acceptable as players often take irrational decisions and the ANN only tries to find an approximation of the results.

We use an iterative development method to design the ANN. We start by designing a simple ANN and then move on to more complex ANN's.

All our ANN's have exactly two outputs. The first output is whether or not to be defensive by checking or calling and the second output is whether or not to be aggressive by betting or raising. The closer an output is to the value one, the more certain the ANN is, that it is the correct decision.

All our ANN's use normalised inputs and a sigmoid function as transfer function. The reason we settle for a sigmoid function rather than a step function is because the sigmoid function allows us to see how certain the ANN is of each decision being correct.

2.1.1 Artificial neural network (ANN)

An artificial neural network (ANN) is inspired by the human brain. It can be used for pattern recognition or classification among other things. An ANN can take any number of inputs and return any number of outputs.

An ANN is made up of neurons that are connected into a network. Each neuron takes a set of inputs and returns a single output. The output of a neuron is sent to all the connected neurons. Each input has a weight that determines influence of the input. The neuron uses an input function to calculate the net input, usually the sum of all weighted inputs, and pass

it on to the transfer function. The type of transfer function determines the output. A step function returns zero or one if the net input is above a certain threshold. This is useful for logical functions. If one needs a value between zero and one a sigmoidal function can be used instead.

Figure 4 models a neuron. Here the weighted inputs w_{i1} , w_{i2} , and w_{i3} are all sent to the input function Σ which then calculates the net input net_i . The transfer function f then calculates the output from the net input.




Figure 4: Model of a neuron.[10]

The neurons in an ANN are distributed in layers as shown in figure 5 and 6. The coloured circles represents neurons and the arrows represents the connections between the neurons. There are three layers, an input layer, a hidden layer, and an output layer. An ANN consists of one input layer and one output layer but may contain any number of hidden layers. The simplest type of ANN is the perceptron which has no hidden layers, see figure 5. It is used for single calculations. A multilayer perceptron is another type of ANN which contains hidden layers. It is used for more complex domains with multiple layers of computations.




Figure 5: Perceptron with two inputs and one output.




Figure 6: Multilayer perceptron with three inputs and two outputs.

The ANN can be trained using a training set of inputs. Using supervised learning, in contrast to unsupervised learning, one must also supply an expected output. For each input it will adjust the weights in order to get closer to the expected output. The TNE indicates the amount of training data that did not produce the expected result. The TNE is calculated during each iteration and the ANN will continue adjusting the weights until the TNE is acceptable.

Backpropagation is the most common algorithm for supervised ANNs. It adjust the weights from the end (the output neurons) back to the start (the input neurons).

After the training the ANN can be validated to see if it works. This is done using a test set different from the training set and see if the results of the ANN matches the expected results of the test set.

2.1.2 First ANN design

For the first attempt we design a simple perceptron that takes two inputs, and returns two outputs, see figure 7.




Figure 7: First ANN design.

The first input is the hand strength. We use the calculator from chapter 1 to calculate the probability of winning against a single opponent. The reason we always find the probability against a single opponent is because the probability of winning decreases drastically as the number of opponents increases. We use an absolute hand strength rather than a hand strength relative to the number of players. This makes it easier to compare the hands strengths in situations with different numbers of players.

The second input is the number of opponents. The input is normalized as

$$I_{norm} = \frac{Opp}{Opp_{max}}$$

Here Opp is the number of opponents and Opp_{max} is the maximum number of opponents (in our case nine).

From the tests we found that this network had a TNE of $\sim 19,9\%$, see section 2.2.1. This does not fulfil our requirement of a TNE of five percent or less.

In order for a perceptron to be accurate the data have to be linear separable. This means that the outcomes can be separated by a single line if they are plotted in a graph. We plot some of the data in figure 8 to see if it is linear separable. The figure clearly shows that the data is not linear separable and therefore a single perceptron will not work.




Figure 8: Distribution of actions. The x-axis the number of opponents and the y-axis is the hand strength. Aggressive actions are indicated by red dots and defensive actions by blue dots.

2.1.3 Second ANN design

Since the first approach designing a perceptron did not fulfil our requirement we instead design a multilayer perceptron (MLP) with five inputs, see figure 9.

The MLP is taking the same inputs as the perceptron from section 2.1.2, but now it takes three additional inputs: The chips of the player, the cost for the player to call, and the pot. We normalize the three new inputs as:

$$I_{norm} = \frac{I}{chips_{total}}$$

. Here I is the input to be normalized and $chips_{total}$ is the total amount of chips in the game, including the pot and the bets.

The MLP has one hidden layer with two hidden neurons. One hidden neuron to calculate the likelihood of winning and another for the economically aspect.




Figure 9: Second ANN design.

From the test in section 2.2.2, we can see that the TNE is $\sim 19,1$ %. This is a small improvement compared to the first ANN but it still does not fulfil the requirement about a TNE of five percent or less.

2.1.4 Third ANN design

For our first ANN we use the same structure as the second ANN, see figure 9. We believe that the reason the second ANN did not work is because the data is found of random players. Each player may have a different strategy resulting in different actions. This will make it hard for an ANN to find a pattern.

Instead we use data of a single player in all games the player is relevant. The test in section 2.2.3

2.2 Test

The University of Alberta has a research group that specializes in the field of artificial intelligence in poker [9]. They have released a dataset containing data from ~ 18.000 real-life rounds of Texas hold'em limit poker. The dataset only contains data about the hole cards of the players who made it to the showdown. We will refer to these players as relevant players.

The dataset consists of multiple files. The file called *hroster* contains the names of all players at the beginning of the round. The file called *hdb* contains the information about the community cards and the flop after each

game state. For each player a file called *pdb.<playername>* exists. This file includes the hole cards (if the player did not fold), chips, profit from round and actions during each game state.

We have collected all the data from the different files into a single data file called *refactored-data.txt* to make it easier to analyse the rounds. We discarded the data for all rounds that has no relevant players. This is because no information about the hole cards exists in such rounds and therefore it is irrelevant for us.

refactored-data.txt will be used for training and testing the ANN's throughout this section.

For each test we create a new dataset. This dataset contains the information about the game for each action performed by a relevant player. All informations about the game is found in the moment of the action. For instance, an action in the pre-flop will have no information about the community cards. Each dataset contains data from 200 random poker rounds, resulting in ~ 1200 actions distributed across all game states.

To create and test the ANN's we use a framework called Neuroph [11]. Neuroph is a neural network framework programmed in java.

2.2.1 Find the TNE of the first ANN

For this test the dataset only includes the normalized data about the hole cards of the player, the community cards and the number of opponents.

To find the TNE of the perceptron described in section 2.1.2, we train it using the dataset. The graph seen in figure 10 displays the TNE through each iteration of the training. From the graph we can see, that the TNE does not get any smaller than $\sim 19,9$ %.




Figure 10: TNE graph from training the first ANN.

2.2.2 Find the TNE of the second ANN

For this test the dataset also includes the normalised data about chips, cost and pot.

We train the MLP from 2.1.3 and get the NTE graph shown in figure 11. The TNE does not get any lower than $\sim 19,1$ %.




Figure 11: TNE graph from training the second ANN.

2.2.3 Find the TNE of the third ANN

For this test the dataset contains the same data as the previous test but all data is found for a single player.

We train the MLP from 2.1.4 and get the NTE graph shown in figure 12. The TNE does not get any lower than $\sim 19,1$ %.




Figure 12: TNE graph from training the third ANN.

2.3 Discussion

When artificial intelligence starts out by playing the games of poker it doesn't have any kind of information about the opponents. So to make sure that the bot won't just go keep throwing away the bank roll. Instead we wanted the bot to minimize our losses so that we would still have a decent bank roll when we have gathered information about the opponents so that the bot could make qualified guesses at what move would be the most appropriate in terms of the current opponent. The default player was never meant to be on the same level as a human player, but humans are only in a slightly better position than the default bot. The human player can like the bot only see the hole and community cards, but a human is also able to make a profile of the bot in their head. This means that they can learn how the bot decides and exploit this. This is one of the reasons that we chose to go with a neural network. In a neural network we are able to of course train the network to make correct decisions based on the targeted output that we give it. But as the game proceeds the neural network has an input which will weight when we have enough information about the opponents to shift our gameplay from the default play style to a more adaptive one. The time that it takes a human

player to learn about the bot and adapt to it, should be the same for the bot. So if we imagine a human player who is good enough to decipher the way the bot is playing and adapt to it. When the human player does that, the bot should also have started to change its ways. Slowly as the bot learn it will adapt more and more to the opponent so when we have enough information the bot will shift completely and disregard the default play style.

2.4 Conclusion

3 Learning to adapt to opponents

In the previous chapter we created an algorithm to develop the default strategy for the APC. The default strategy is one of the two strategies that will be used to determine the actions of the APC. The other strategy is the adaptive strategy.

The goal of this chapter is to learn the APC to create an algorithm to develop the adaptive strategy by solving problem statement 3.

Problem statement 3

How can one develop a strategy adapted to the opponents?

In poker there is no such thing as a single optimal strategy. Every strategy has weaknesses and therefore the optimal strategy is one that takes advantage of the weaknesses of the strategies of the opponents. In order to take advantage of the opponents strategies one must first understand their strategy. In poker understanding the opponents is one of the most important key elements of the game.

Once one understand the strategy of the opponent one has to adapt one-self's strategy.

3.0.1 Player modeling

Player modeling is a loosely defined concept and may vary from one context to another. The concept of player modeling is to make a computational model of a player. This model includes game related attributes, such as play style and preferences, as well as non-game related attributes, such as cultural background, gender, and personality. All decisions of the player are ultimately made on the basis of these attributes.

Player modeling is used to describe or predict the players decisions, reasoning and reactions. In the field of artificial intelligences the human player is the most used model for developing computer players. Understanding the reason behind every choice of a player will not only bring a better understanding of the player but also a better understanding of the game and its mechanics.

Since the player model can easily become extremely extensive one normally only includes the relevant attributes of the player.

If player can optimally model an opponent, that player will have a much greater chance of being successful. A player who have the same amount of knowledge regarding the fundamentals of poker as the opponent, but lacks the ability to change strategy throughout the game will probably not win as much as the opponent.

3.1 How can we model a player?

Quite a few decisions needs to be made when beginning to player model. It is crucial to figure out which attributes are relevant to create the specific type of player model that one wishes to implement. The attributes we find relevant is the following:

1. The opponents aggressiveness
2. The opponents current aggressiveness
3. The opponents bluff percentage
4. The opponents reaction time

Each of these attributes and the reasoning behind the choices is described below.

As being stated earlier, poker is a game of imperfect information. The betting strategy of a player is the only piece of information that an opponent is able to see. But just because a player is choosing to raise, does not necessarily mean that the player has a hand that has a good chance of winning. A player who knows the fundamentals of poker are aware that the way of betting could reflect what cards he is holding. A player will therefore often try to bet the opposite of what is logic in order to disguise his cards further. It goes without saying that predicting the cards of an opponent is rather impossible. This is why we chose aggressiveness as one of the attributes for the player modeling. The opponents aggressiveness can be calculated by looking at the history of the game. If an opponent raises in every game the system will address him as an aggressive player and calculate a value describing how aggressive he is. Likewise if the opponent is limiting his actions to bet/check then the aggressiveness value will decrease.

The opponents current aggressiveness has been chosen in order to make APC prepared for any changes in the opponents characteristics. If an opponent has been playing safe the whole game and suddenly begins to raise, it is quite possible that he is either changing strategy or maybe he has received some good hole cards. If the opponent has a habit of raising each time he receives some good hole cards, it says a lot about his current strategy. However, the main focus will still be in regards to the player changing his strategy which APC will have to adapt to.

The opponents bluff percentage is chosen as a relevant factor in regards to a modeling of the opponent as it tells whether or not the opponent is playing solely from his point of view. If an opponent is playing solely from his own point of view and not thinking about the see-through he creates when doing this, then APC will have a great advantage. After several games, APC will notice a pattern for when the opponent is raising compared to the cards shown at the showdown. APC will then for the further games against this opponent be able to distinguish between the opponent bluffing or not.

The opponents reaction time may or may not be relevant when modeling the opponent. The reaction time used by the opponent when making a decision does not necessarily mean that the opponent is uncertain if he should do this or that. Many factors come into play when an opponent is having his turn. A delay in reaction can be due to the opponent being away from the keyboard (AFK), whether or not he should fold, bet or raise based on his hole cards, the community cards, APCs action among other things.

3.2 How can we model a player dynamically?

In order to make the player modeling dynamic we will have to introduce an ANN. The ANN can learn dynamically based on the inputs given. The network keeps getting new inputs that can be used in the current dataset for an opponent and should be added so that the network can learn each time the opponent makes a decision that gives us more data.

3.3 Test

3.4 Discussion

3.5 Conclusion

4 Discussion

5 Conclusion

Glossary

Poker

Texas Hold'em

The most popular variance of poker. Game rules [7]

Limit / No-limit

Restrictions about how much players are allowed to bet or raise. No-limit has no restrictions while in limit you are only allowed to bet and raise a fixed amount.

Poker game

A game is from the players start playing to all except one player have been knocked out or left the game.

Poker round

A round is from the poker get their hole cards dealt to the winners are found. In each round the players will have new hole cards.

Game state

A poker round is divided into four game states: Pre-flop, flop, turn, river. The difference between the states are the number of community cards dealt. The numbers of community cards dealt are zero, three, four, five for each state respectively.

Pre-flop

The first game state. Each player is dealt two cards and zero community cards are dealt. After the bidding round the flop will be dealt.

Flop

The second game state. In this state three community cards are dealt. After the bidding round the turn will be dealt.

Turn

The third game state. In this state one community card is dealt making it a total of four community cards. After the bidding round the river will be dealt.

River

The fourth and final game state. In this state one community card is

dealt making it a total of five community cards. After the bidding round the winner of the round will be found.

Preflop: Preflop is the stage before the flop. Each player has been giving 2 cards each.

Flop: Flop is the stage when the flop cards are on the table (3 cards)

Turn: Turn is the stage when the turn card is on the table (1 card extra so 4 cards in total)

River: River is the stage when the river card is on the table (1 card extra so 5 cards in total) Game: A full game is when one of the players has lost and by lost we mean hitting 0 in the bankroll. Round: A round is when we have been through all stages, preflop, flop, turn, river and the end of the round when one of the players has won the round by having the best cards.

Hand: Hand is the 2 two cards a player is holding, therefore each player has a hand. Pot: Pot is the sum of the total amount each player has betted in the given round. Deck: Deck is the deck of cards Bluff: Bluff is when a player is trying to make the opponent think that he has good cards on his hand by raising or calling when actually he has nothing or very bad cards. Aggressive: Aggressive is when a player or computer is playing aggressively Defensive: Defensive is when a player or computer is playing aggressively All-in: All-in is when a player is betting all of his bankroll Bankroll: Bankroll is the total amount that a player has to bet with, when this hits 0 the player has lost the game. Board: Board is the table that we are sitting and playing at. Small blind: Small blind is an amount that the player who has the small blind HAS to pay even though they want to fold their hand. Small blind is half the amount of big blind. Big blind: Big blind is an amount the player who has the small blind HAS to pay even though they want to fold their hand. Big blind is double the amount of small blind. Check: Check is when a player doesn't want to bet but just wants the game to go on. Call: Call is when the opponent has betted an amount and if we want to proceed the round we will have to call his bet and lay the same amount into the pot as he did. Raise: Raise is when the opponent has betted an amount and if we want to proceed with the round, we can either call his bet or raise the amount by laying more than what our opponent laid into the pot. Thereby our opponent will have to either call, fold or raise. Fold: Fold is if you don't want to play with the cards you have been dealt or the opponents bet is too high for you, given the cards you have on your hand. Limit: Limit is so that one playing cannot just bet 100 kr. Or whatever we set the limit to. But perhaps the limit is 10kr

for each bet/raise. No-limit: No-limit means that a player can bet whatever he wants or raise with whatever amount he has left in his bankroll. Heads up: In a heads up we are only having 2 players who play against each other.

References

- [1] caniwinn. (no date known) Texas Holdem Heads-Up Preflop Odds [Online]. Available: <https://caniwin.com/texasholdem/preflop/heads-up.php>
- [2] inferisx. (no date known) Top 10 popular card games [Online]. Available: <http://topyaps.com/top-10-popular-card-games>
- [3] Marlos C. Machado, Eduardo P. C. Fantini and Luiz Chaimowicz *Player Modeling: What is it? How to do it?*, SBC - Proceedings of SBGames 2011.
- [4] Aaron Davidson, Darse Billings, Jonathan Schaeffer and Duane Szafron, *Improved Opponent Modeling in Poker*, Department of Computing Science, University of Alberta.
- [5] Garrett Nicolai and Robert Hilderman, *Algorithms for evolving no-limit Texas Hold'em poker playing agents*, Department of Computer Science, University of Regina, Regina, and Dalhousie University, Halifax, Canada.
- [6] Jochen Fröhlich, *Neural Net Components in an Object Oriented Class Structure*, ebook pp. 11-28
- [7] Poker Listings, <http://www.pokerlistings.com/poker-rules-texas-holdem>
- [8] Lily Hay Newman, *Using AI to Study Poker Is Really About Solving Some of the World's Biggest Problems*, Slat blog
- [9] University of Alberta Computer Poker Research Group, <http://poker.cs.ualberta.ca/>
- [10] Artificial neuron models, <http://www.willamette.edu/gorr/classes/cs449/ann-overview.html>
- [11] Java Neural Network Framework, <http://neuroph.sourceforge.net/>
<http://neuroph.sourceforge.net/>

A

Glossary